



Architecture of reconfigurable systems : Partial reconfiguration mechanisms

J. Lorandel

jordane.lorandel@u-cergy.fr

Slides available there: https://perso-etis.ensea.fr//lorandel/M2_ESI_ASR.php

□Outline

1. Resources
2. Partial Reconfiguration
 - Reconfiguration modes in FPGA
 - Why use (dynamic) partial reconfiguration?
3. Methodology
 - Different ways to perform PR
 - Module-based partial reconfiguration
 - Difference-based partial reconfiguration
4. Conclusion

Resources

- This course uses a Xilinx FPGA board. There are plenty of resources available out there (mostly online)
- From Xilinx Inc.
 - Xilinx Zynq-7000
 - <http://www.xilinx.com/products/zynq-7000/third-party-documentation.htm>
 - See : http://www.xilinx.com/support/documentation/data_sheets/ds180_7series_Overview.pdf
 - Reference Manual : Zynq-7000 All Programmable Technical Reference Manual **UG585**

Partial Reconfiguration – Xilinx

- Overall Xilinx documentation page for partial reconfiguration:

<https://www.xilinx.com/products/design-tools/vivado/implementation/dynamic-function-exchange.html#documentation>

- See also :

➤ User Guide:

https://www.xilinx.com/support/documentation/sw_manuals/xilinx2017_1/ug909-vivado-partial-reconfiguration.pdf

➤ Tutorial: https://www.xilinx.com/support/documentation/sw_manuals/xilinx2017_4/ug947-vivado-partial-reconfiguration-tutorial.pdf

Partial Reconfiguration – Altera

- Overall page:

<https://www.intel.com/content/www/us/en/programmable/products/design-software/fpga-design/quartus-prime/features/partial-reconfiguration.html>

- See also:

- User Guide: <https://www.intel.com/content/www/us/en/programmable/products/design-software/fpga-design/quartus-prime/user-guides.html>
- Tutorial: <https://github.com/intel/fpga-partial-reconfig/tree/master/tutorials>
- Intel PR page: <https://01.org/fpga-partial-reconfiguration>

□ Partial Reconfiguration

WHY ??

□ Lack of hardware resources

- Not enough room to include the logic of all possible hardware functions in the target
- Need to reduce the size of the resulting device

□ Flexibility

- Implementation of various features with different algorithms
- Implementation of various protocols to achieve the same end-goal

WHY ??

❑ Reliability and security

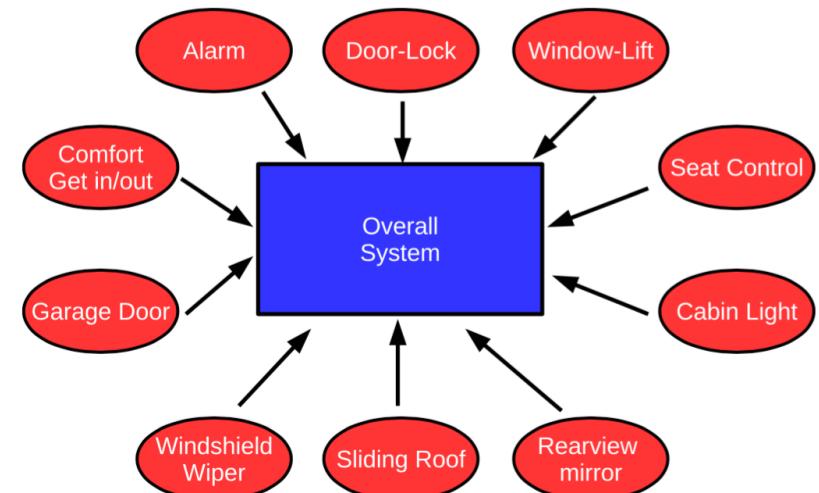
- Ensure the overall system can withstand faults due to external factors
 - Bit flips due to solar flares, radiation, etc.
- Guarantee some circuits functions are only available to specific users/devices
- Adapt cryptographic protocols to context

❑ Increase reconfiguration speed when needed

- i.e., instead of stopping the whole circuit before reconfiguring it

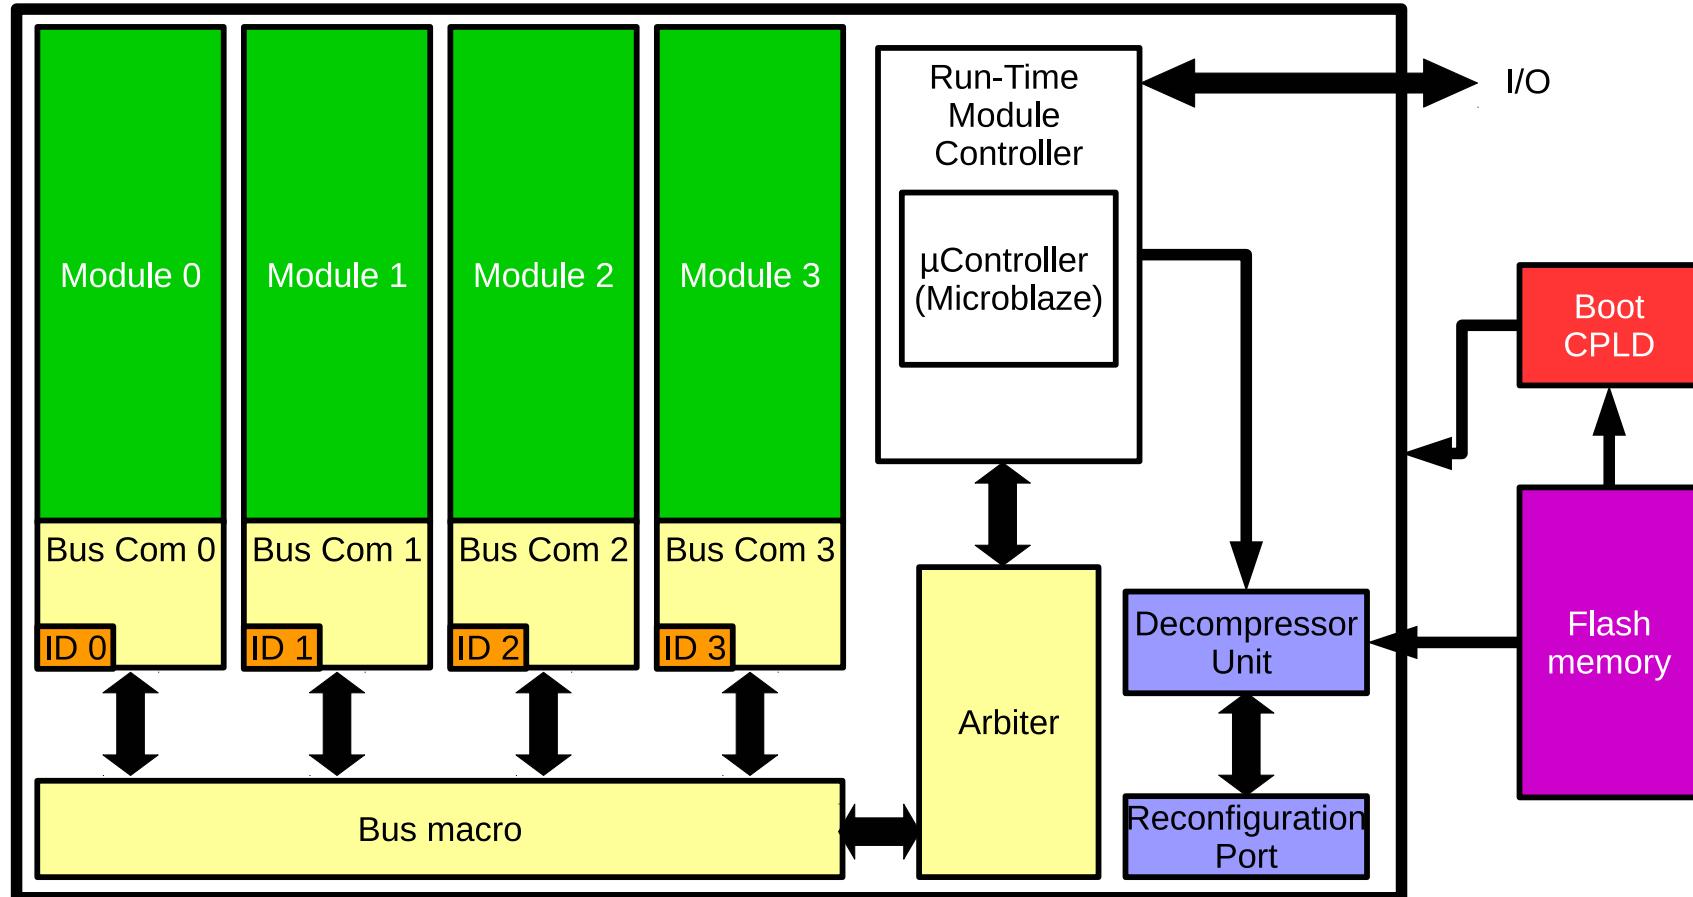
□ Some examples

- Cognitive radio
 - Adapt the modulation and coding scheme and type depending on the environment; [Rihani16]
 - Modulation, FEC (Reed-Solomon, turbo-code,...), data rate
 - Adapt the wireless communication standard and resource allocation ?
- Automotive industry [BeckerEtAl07]
 - Reduce the amount of logic embed in a vehicle by reconfiguring HW task (depending on criticality)



Some of the cabin functions of a modern vehicle

WHY ??



Set of tasks to perform:

- **Communication management**
 - Delivering messages to/from the functions
 - Storing messages
- **Reconfig. Management – context switching**
 - Finding available slot
 - Saving current state
 - Sending instruction block
 - Restoring state
- **Resource Management**
 - Bookkeeping of busy/idle modules
 - Management of message buffers for each function
 - Storing state variables

WHY ??

Let's play with your imagination

- > Autonomous vehicle, personalization?
- > IoT object with support of multiple wireless communications standards?
- > Energy-efficient reconfigurable devices?
- ...

❑ Static Reconfiguration

- Load the whole bitstream on the FPGA
- Stop the execution to do the loading

❑ Static Partial Reconfiguration

- Only a portion of the bitstream is load
- Stop the execution to load the partial bistream
 - But the wait time is much shorter

❑ Dynamic Partial Reconfiguration (DPR)

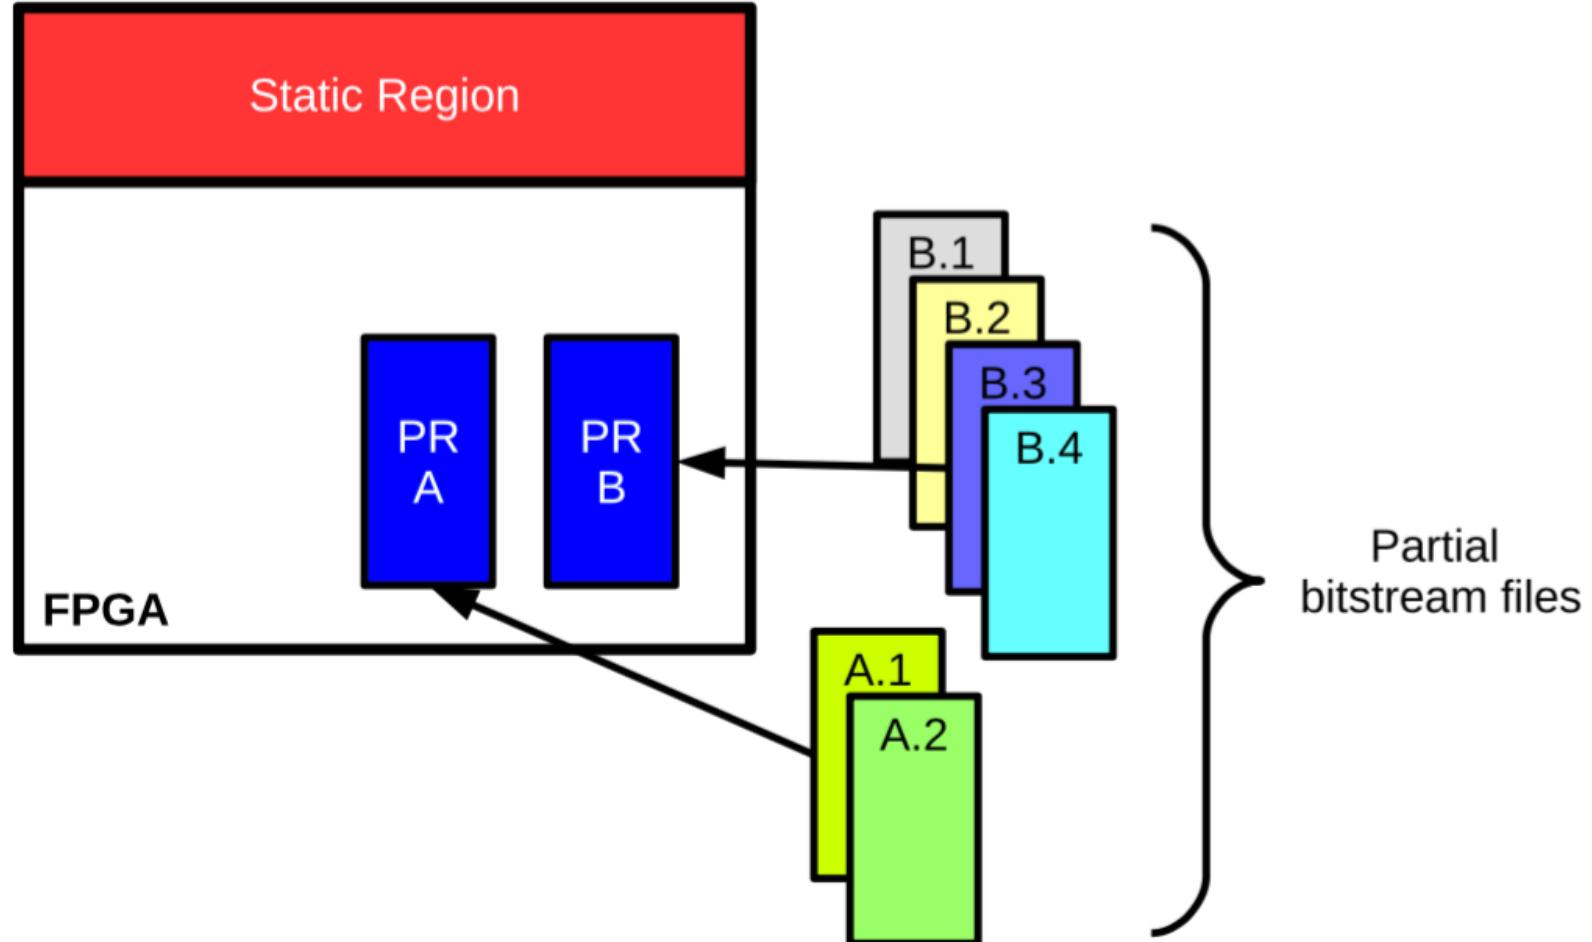
- Part of the FPGA is being reconfigured while the rest is still running

Configuration types

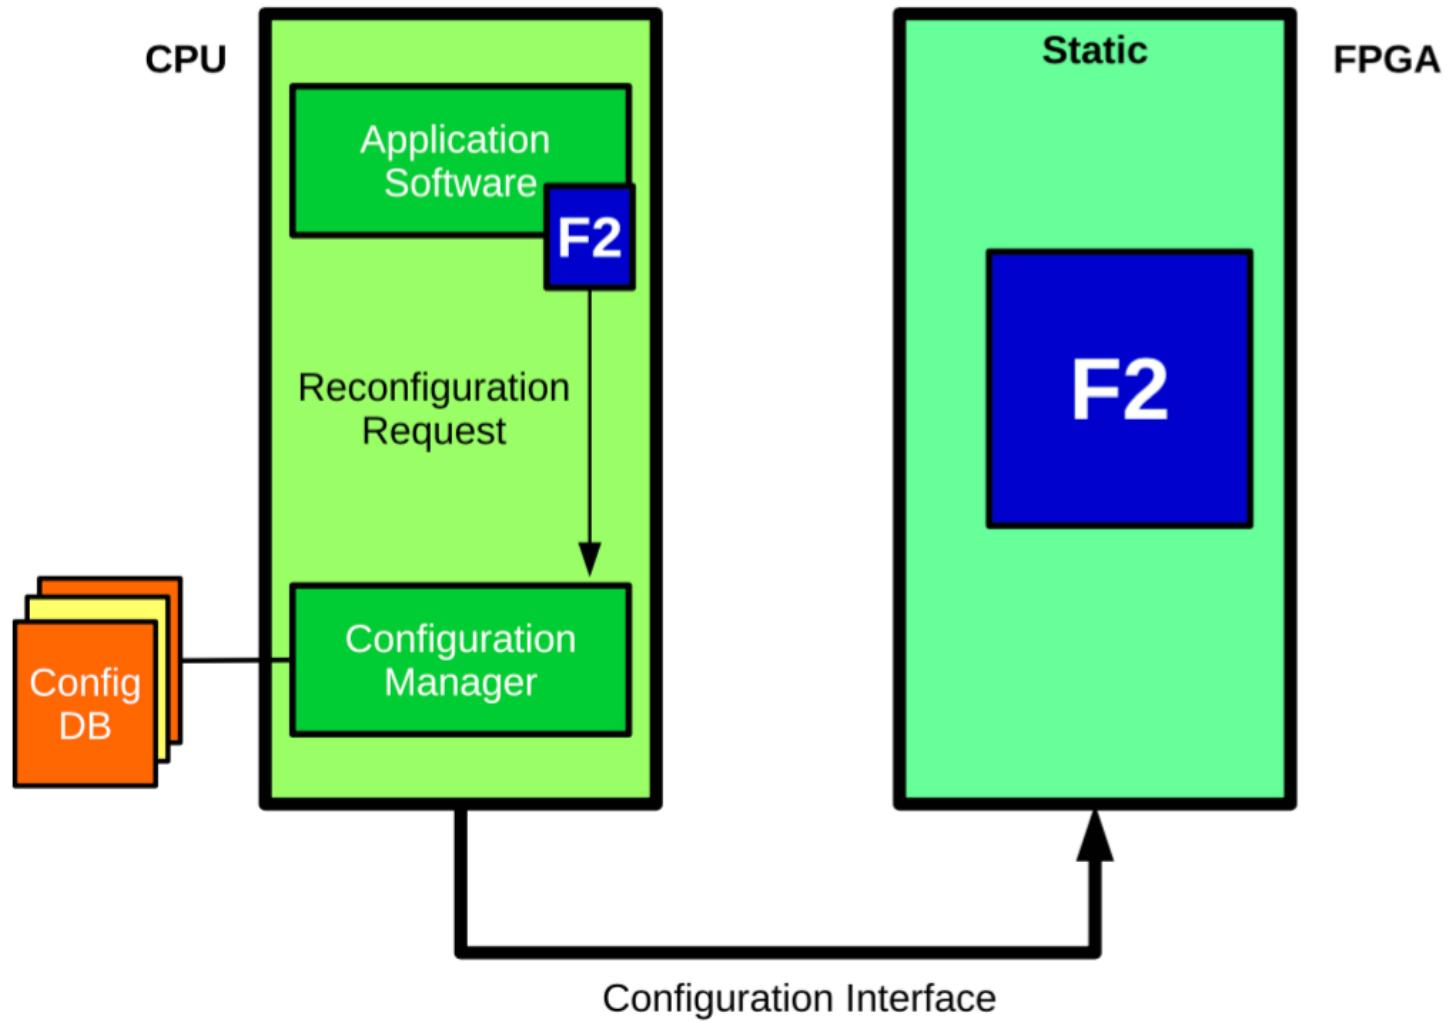
This course tries to provide general concept but some mechanisms are vendor-specific

- We'll mostly discuss Xilinx FPGAs for PR
 - Xilinx is one of the first vendors to propose mechanisms for (D)PR
- Altera FPGAs follow a similar path, but their toolchain and capabilities are much more recent than Xilinx'
 - PR features are mostly available for high-end FPGAs
 - There is a clear will of Intel/Altera to make it more accessible/available to everyone
 - They are not there yet

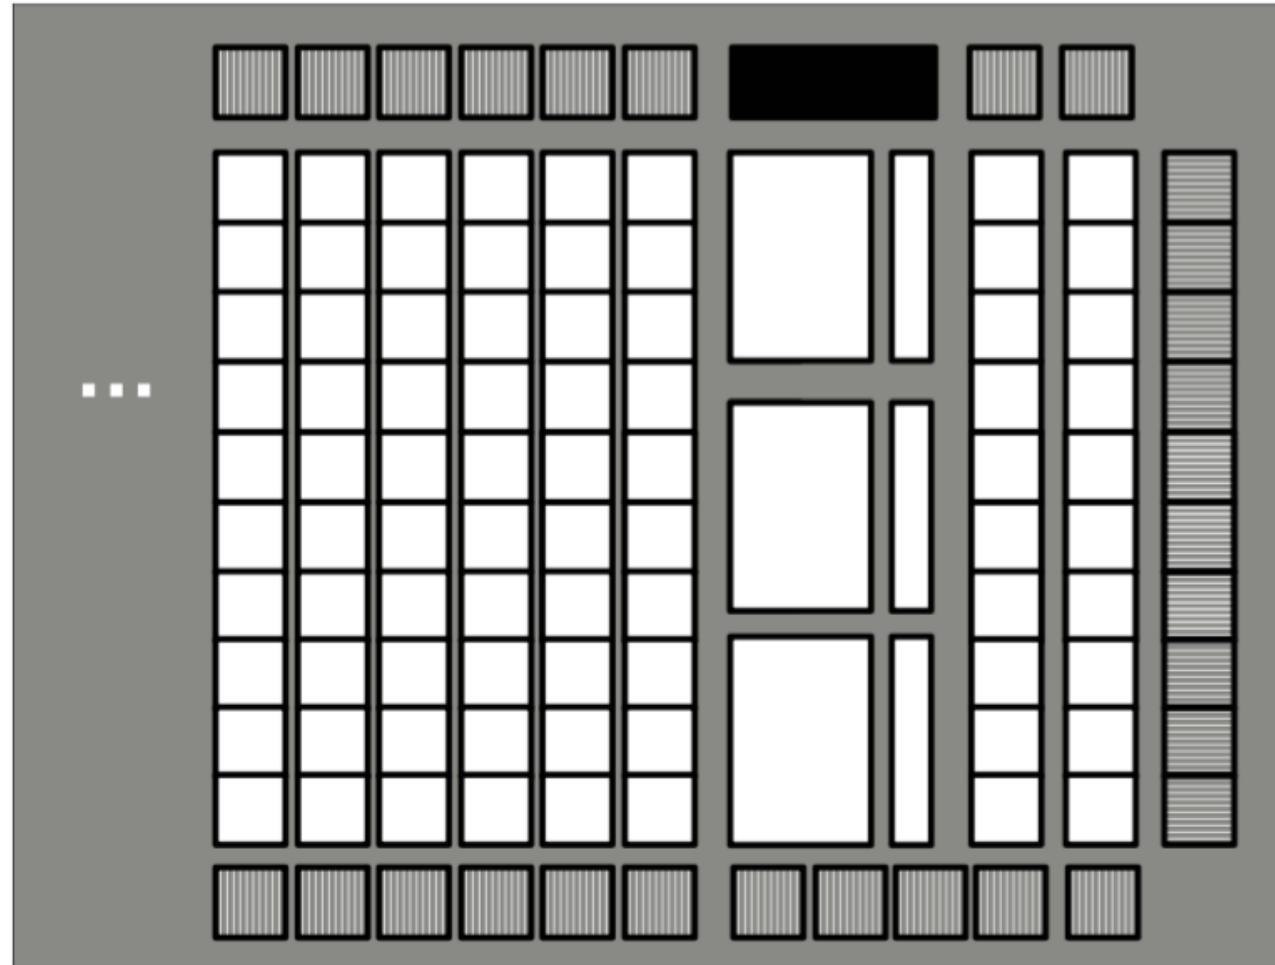
Partial Reconfiguration I



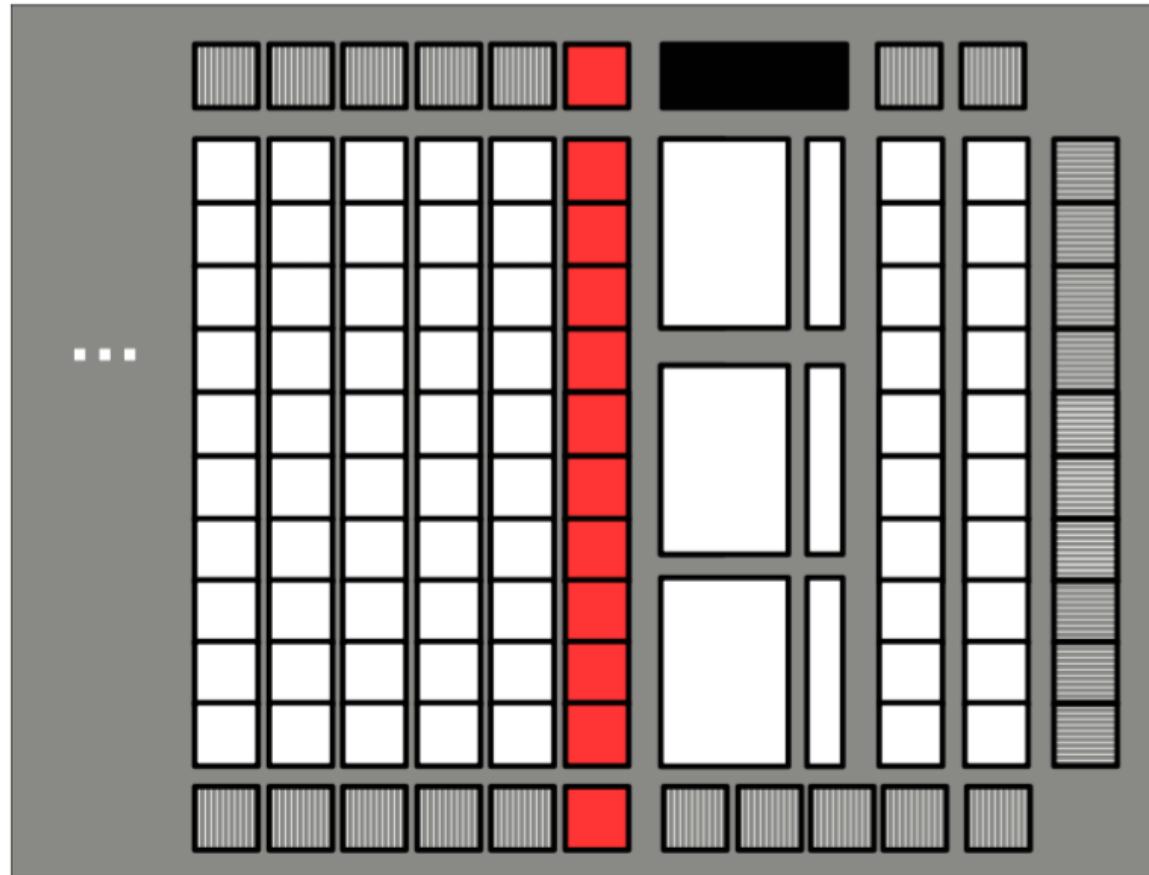
Partial Reconfiguration II



Partial Reconfiguration III

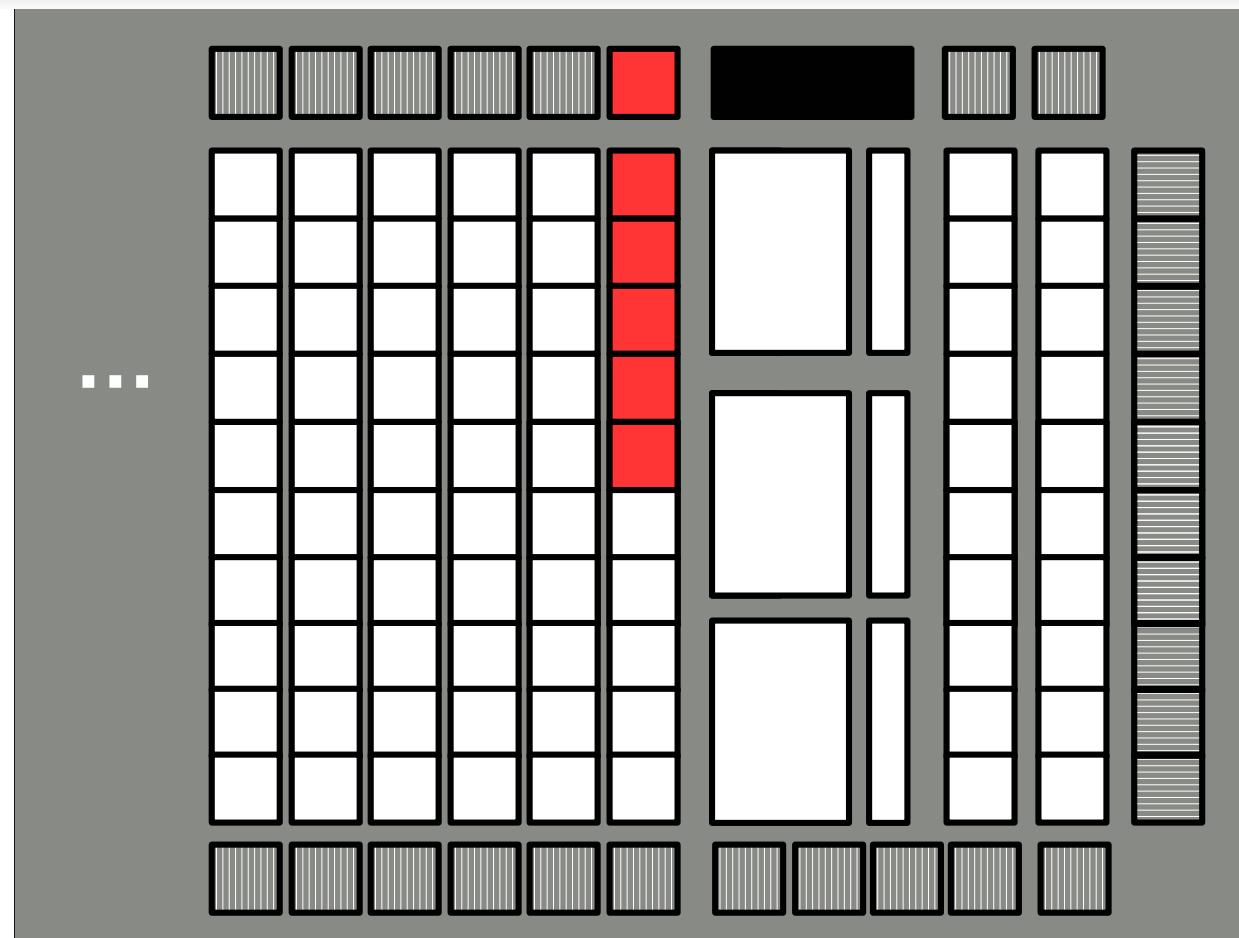


Partial Reconfiguration IV



For older generation : the whole column must be selected to create a partial reconfiguration module

Partial Reconfiguration V



Newer generations of FPGAs allow to use a rectangular area, with a fixed minimal number of rows

❑ Reconfiguration speed

- The reconfigurable time is a function of the size and the organization of the partial reconfigurable regions
- Modern FPGAs allow for arbitrary-shaped regions (a.k.a. rectangular shape)

Bitstream

- **Useful documentation:** 7 Series FPGAs Configuration, UG470, p104
https://www.xilinx.com/support/documentation/user_guides/ug470_7Series_Config.pdf
<https://www.kc8apf.net/2018/05/unpacking-xilinx-7-series-bitstreams-part-1/>
- **File formats:**
 - .BIT Binary file containing BIT header followed by raw bitstream
 - .RBT ASCII file with text header followed by raw bitstream written as literal '0' and '1' characters for each bit
 - .BIN Raw bitstream
 - .MCSPROM file format (includes address and checksum info)
- Even though a BIN contains all the necessary data for programming a part, BIT is the default format generated by Vivado's *write_bitstream* command and is what I'll focus on.
- **Bit Header:**
 - Information (design name, build date/time, target part name) usually ignored by the chip. Such information is required by other programming tools (Vivado, openocd)
- **Bitstream file**



- **Configuration packets**

- All data formats are described in 32-bit, big-endian words
- All 7 series FPGA bitstreams commands are executed by reading or writing to the configuration registers
- The FPGA bitstream consists of two packets types: Type 1 & Type 2

Type 1 packet Header (followed by data)

Header Type	Opcode	Register Address	Reserved	Word Count
[31:29]	[28:27]	[26:13]	[12:11]	[10:0]
001	xx	RRRRRRRRxxxx	RR	xxxxxxxxxx

Opcode	Name
00	NOP
01	Read
10	Write
11	Reserved

Only 5 out of 14 register address bits are used in 7 series FPGAs
 ⇒ Define a control register

Name	Read/Write	Address	Description
CRC	Read/Write	00000	CRC Register
FAR	Read/Write	00001	Frame Address Register
FDRI	Write	00010	Frame Data Register, Input Register (write configuration data)
FDRO	Read	00011	Frame Data Register, Output Register (read configuration data)
CMD	Read/Write	00100	Command Register
CTL0	Read/Write	00101	Control Register 0

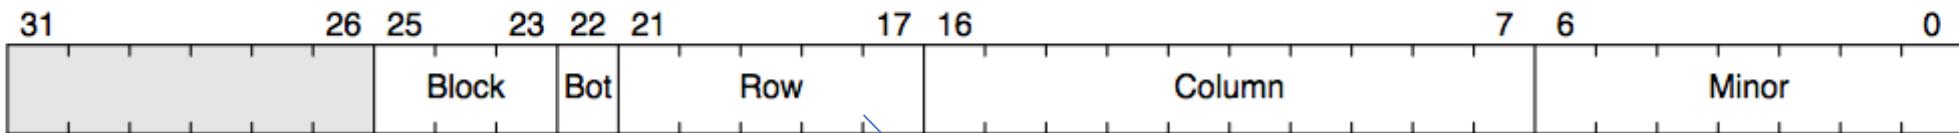
- **Frame Address Register I**

- The 7 series devices are divided into two halves, the top and the bottom. All frames in 7 series devices have a fixed, identical length of 3,232 bits (101 32-bit words).
- The Frame Address Register (FAR) is divided into five fields:

Address Type	Bit Index	Description
Block Type	[25:23]	Valid block types are CLB, I/O, CLK (000), block RAM content (001), and CFG_CLB (010). A normal bitstream does not include type 011.
Top/Bottom Bit	22	Select between top-half rows (0) and bottom-half rows (1).
Row Address	[21:17]	Selects the current row. The row addresses increment from center to top and then reset and increment from center to bottom.
Column Address	[16:7]	Selects a major column, such as a column of CLBs. Column addresses start at 0 on the left and increase to the right.
Minor Address	[6:0]	Selects a frame within a major column.

The address can be written directly or can be auto-incremented at the end of each frame. The typical bitstream starts at address 0 and auto-increments to the final count.

- Frame Address Register II



Xilinx 7-Series Frame Address Register

Code	Block Type
000	CLB,I/O,Clock
001	BRAM content
010	CFG_CLB

Bottom/Up regarding a center-line

Row and column index (0 on left)

Minor frames are contained within a column

- Configuration packets

Type 2 packet Header (must follow a type 1 packet)

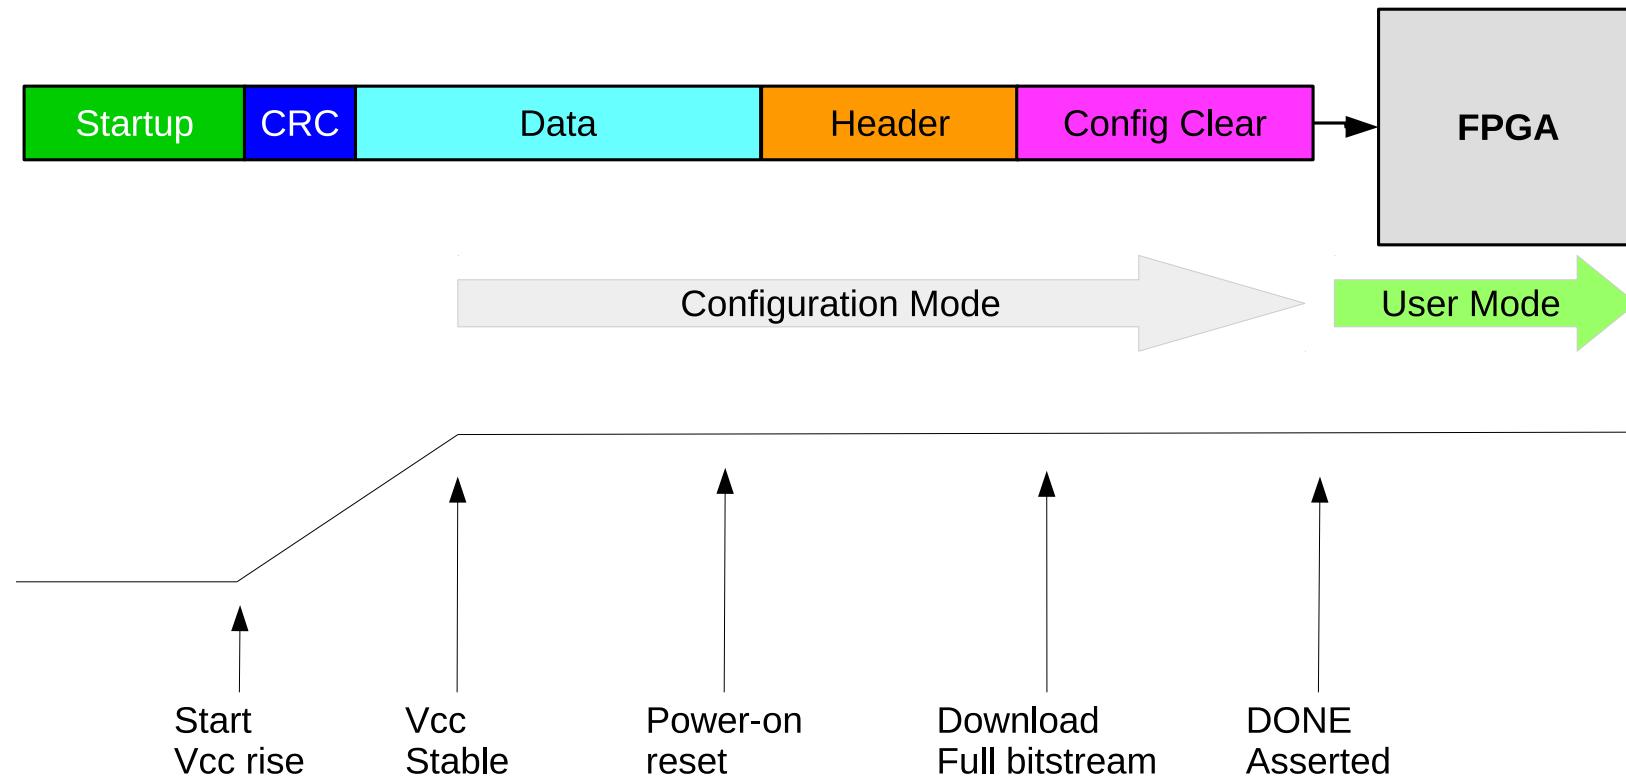
- ⇒ used to write long blocks. No address is presented here because it uses the previous Type 1 packet address. The header section is always a 32-bit word.
- ⇒ Following the Type 2 packet header is the Type 2 Data section, which contains the number of 32-bit words specified by the word count portion of the header

Header Type	Opcode	Word Count
[31:29]	[28:27]	[26:0]
010	XX	xxxxxxxxxxxxxxxxxxxxxxxxxxxx

Number of data word to transmit

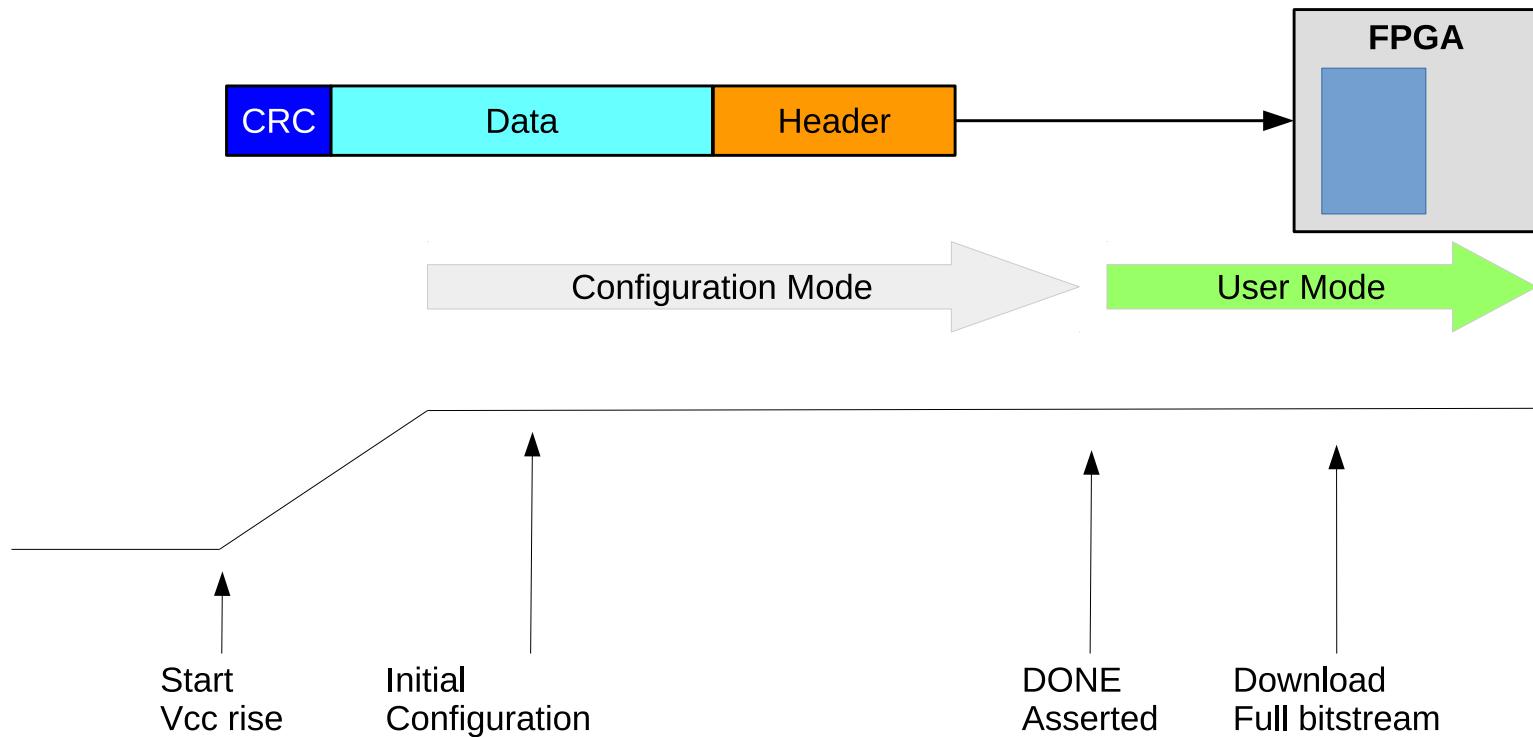
Download Bitstream

- Downloading bit files to FPGA



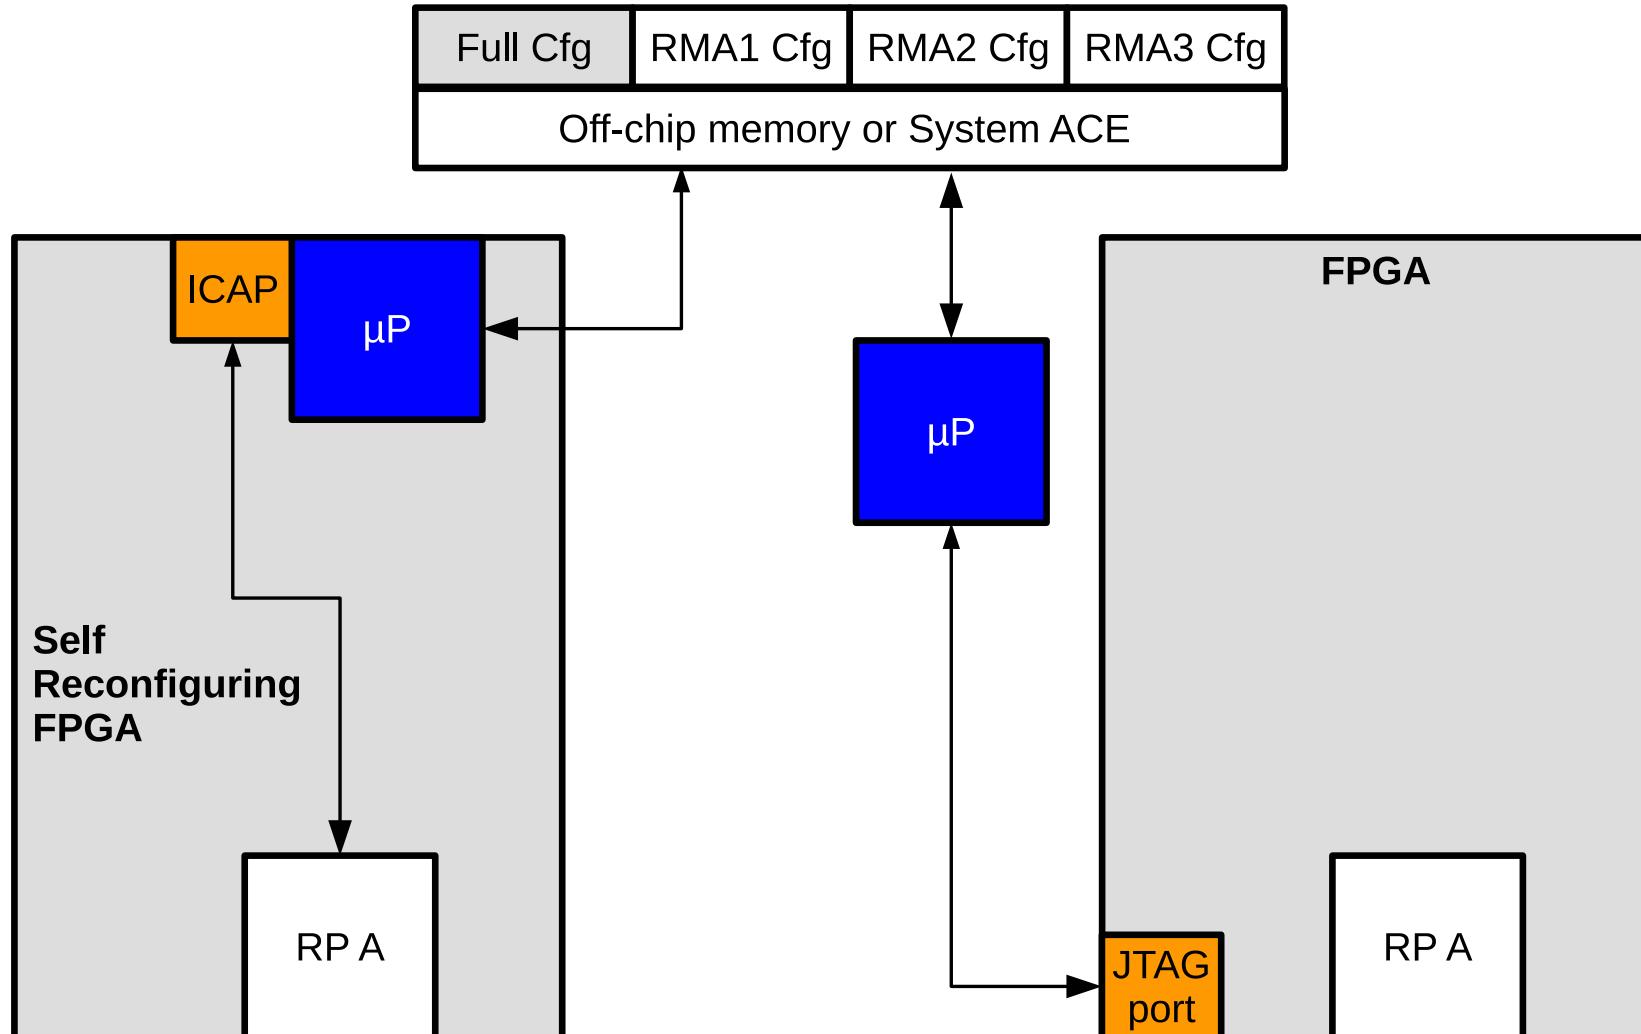
Download Bitstream

- Downloading **partial** bit files to FPGA



- The Partial Reconfiguration control logic can either reside in an **external processor** or in the **fabric of the FPGA** device to be reconfigured.
 - **Externally:** Uses **Serial configuration port**, **JTAG port**, **Processor Configuration Access Port (PCAP)**, or **MCAP***
 - **Internally:** Uses Internal configuration access port (**ICAP**), Internal configuration can consist of either a custom state machine, or an embedded processor such as MicroBlaze™ processor.

PR Overview



□ Two major ways to perform PR

- **Module-based PR**

- Implement reconfigurable modules separately
- Constrain components to be placed at given location
- Complete bitstream built as sum of all partial bitstreams

- **Difference-based PR**

- Implement complete bitstreams separately
- Implement fixed-parts and reconfigurable parts with components constrained at same location within all bitstreams
- Compute the difference of two bitstreams to obtain partial bitstream needed to move from one configuration to the next context

*Focus on this type
in this course*

❑ Component enabling PR

- Basic architectural elements
 - Fixed physical implementation of part of function
 - Stable interface:
 - Inputs specifically targeted at changing functionality
 - Inputs to be manipulated by the component

☐ Configuration frame (Xilinx-specific)

- Smallest addressable segments of FPGA configuration memory space
- 1 frame = {CLB, BRAM, DSP}
 - On more modern/complex FPGAs (UltraScale, etc.), more components may be part of frames. . .
- Bitstream files specify which logic to configure according to frame boundaries

	7-series	UltraScale(+)
CLB	50×1	60×1
DSP48	10×1	24×1
Block RAM	10×1	12×1
IOB / Clock	N/A	$52 I/O \times 1$
Gb transceiver	N/A	4 high

□ Specific components for Xilinx FPGAs

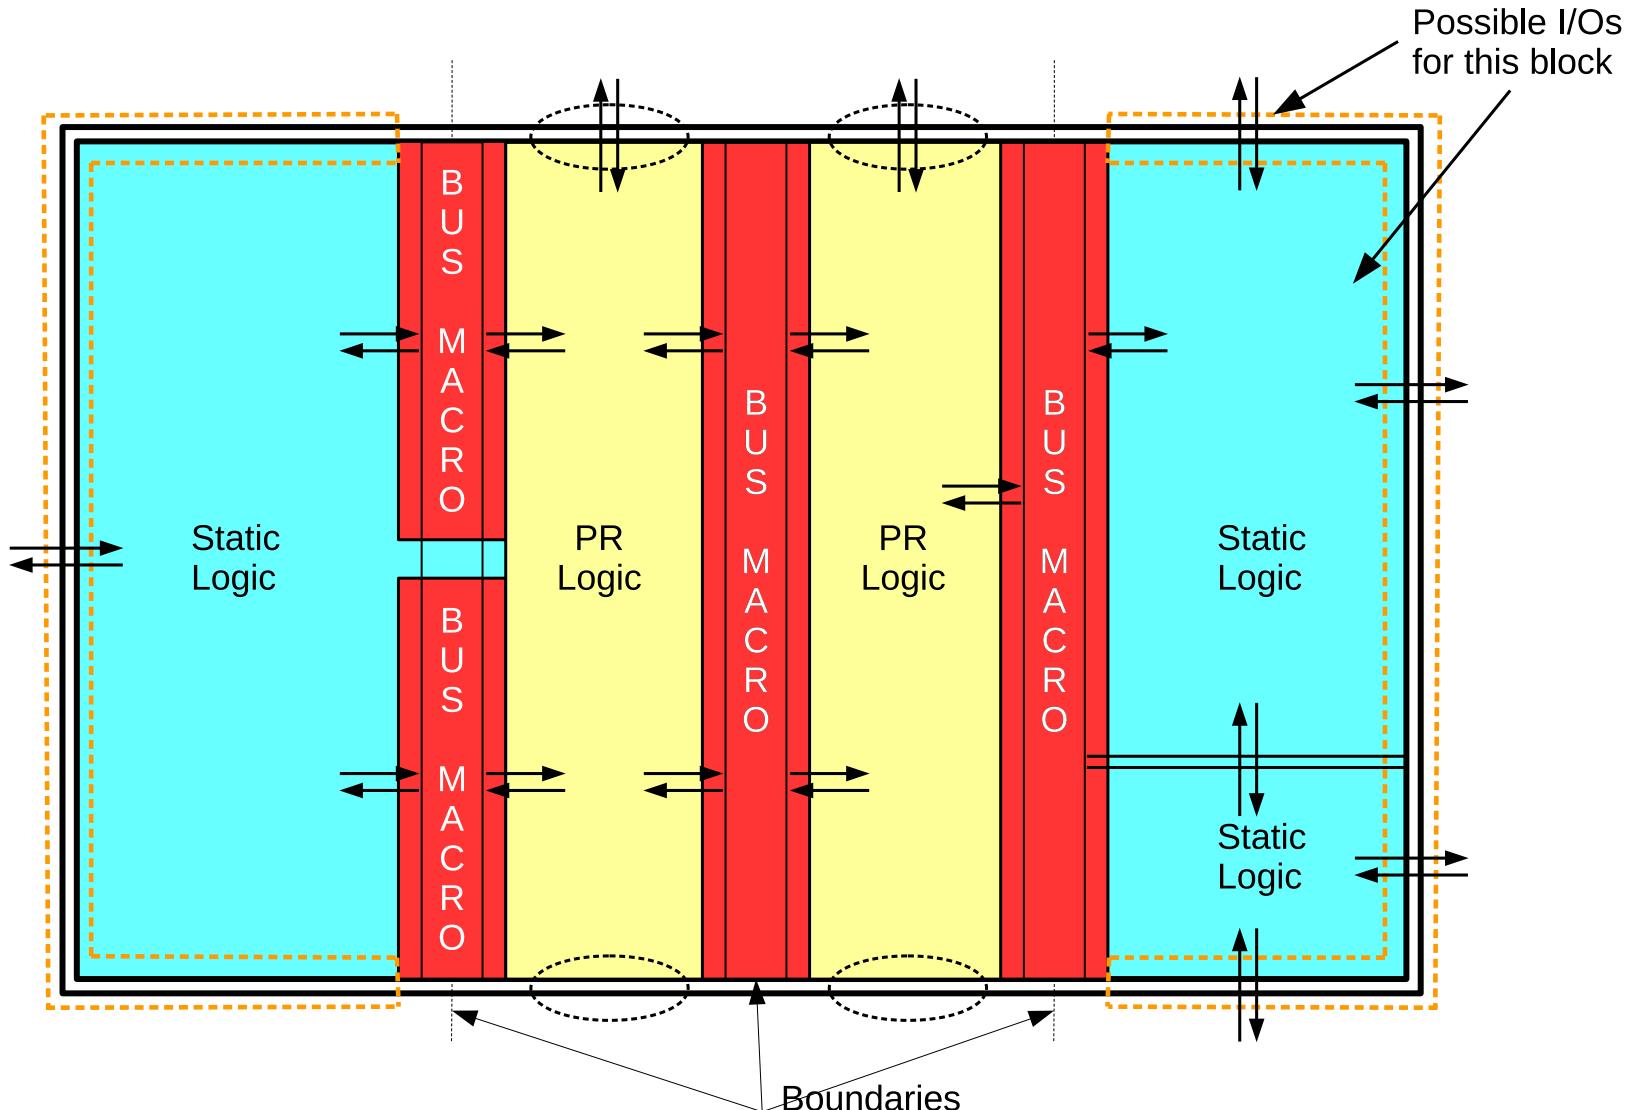
- **Partition** – logical section of the design
 - Defined at hierarchical boundaries
- **Partition Definition (PD)** – set of Partial Reconfiguration Modules (PRM) associated with a module instance
 - A PD is applied to all instances of the module
- **Bus macros** – communication between modules
- **Partition pin** – logical & physical connections between static and PR logic
 - Replace bus macros...
- **Partial reconfigurable module (PRM)** – netlist or HDL description implemented within reconfiguration partition
- **Static logic** – top-level logic, never partially reconfigured
 - Active during PR

□ Specific components for Xilinx FPGAs II

- **Static design** – part of design that doesn't change during PR
 - Includes top-level & all modules not defined as reconfigurable
 - Static design “=“ static logic + static routing
- **Physical block (pblock)** – used to group logic of a reconfigurable module
 - Rectangular area with constraints
 - Motivation: minimize routing between pblocks
 - One pblock should represent less than 20% of the total floor plan

Methodology

□ Bus macro I



❑ Bus macro II

Bus macros provide guaranteed fixed-communication channels among PRMs and static modules

- Designer must ensure signals are routed on the right paths after reconfiguration
- Routing resources must remain static for inter-module communications
- HDL code must be written so that only bus macros are used to have modules communicating with each other
- Each macro: 4-bit inter-communication channel
 - E.G., to enable a 16-bit communication channel, 4 bus macros must be enabled.

❑ Bus macro III -> Partition pins in current FPGAs

New Xilinx FPGAs and tools do not require the designer to explicitly specify bus macros at FPGA boundaries

➤ **Bus macros are replaced by partition pins**

These are the physical interface points between static and reconfigurable logic. They are anchor points within an interconnect tile through which each I/O of the Reconfigurable Module must route (= virtual I/O between RM)

- The tools do the work automatically
- The user can create specific pin partitions when needed
 - Use HD.PARTPIN_RANGE to do so

➤ Modules must still communicate through adjacent IOBs

❑ Physical blocks (pblocks)

- **Snapping mode:** ensures that drawn pblocks are legal (will go to the right row/column boundaries)
- Partitions make it harder to optimize the overall design
 - Don't abuse the use of pblocks
- Reconfigurable frames: pblocks ensure that specific layout constraints are met
- When creating a pblock, one must ensure that
 - **PRM & RP are big enough to include the right amount of BRAM/DSP48 components**
 - **The most complex and/or biggest PRM is designed/outlined first**

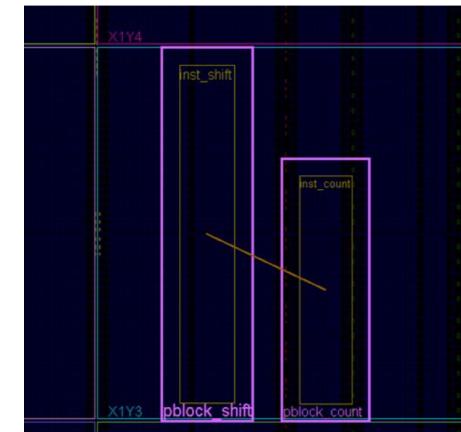
□ Physical blocks (pblocks)

- Reset after reconfig: *specific feature (7 series FPGA) allowing to hold the reconfiguring region in a steady state during PR, and then all the logic of the RM is initialized to its starting value.*

- Very important in synchronous design (FSM, ...)
- The user can specify this pblock feature using :

Set_property RESET_AFTER_RECONFIG true [get_pblocks <reconfig_pblock_name>]

- Pay attention to:
 - Pblock constraints must align to reconfigurable frames (vertically). Because the GSR affects every synchronous element within the region !



Left ok
Right ko

Using the SNAPPING_MODE constraint automatically creates legal, reconfigurable Pblocks

❑ Additional constraints on modular design

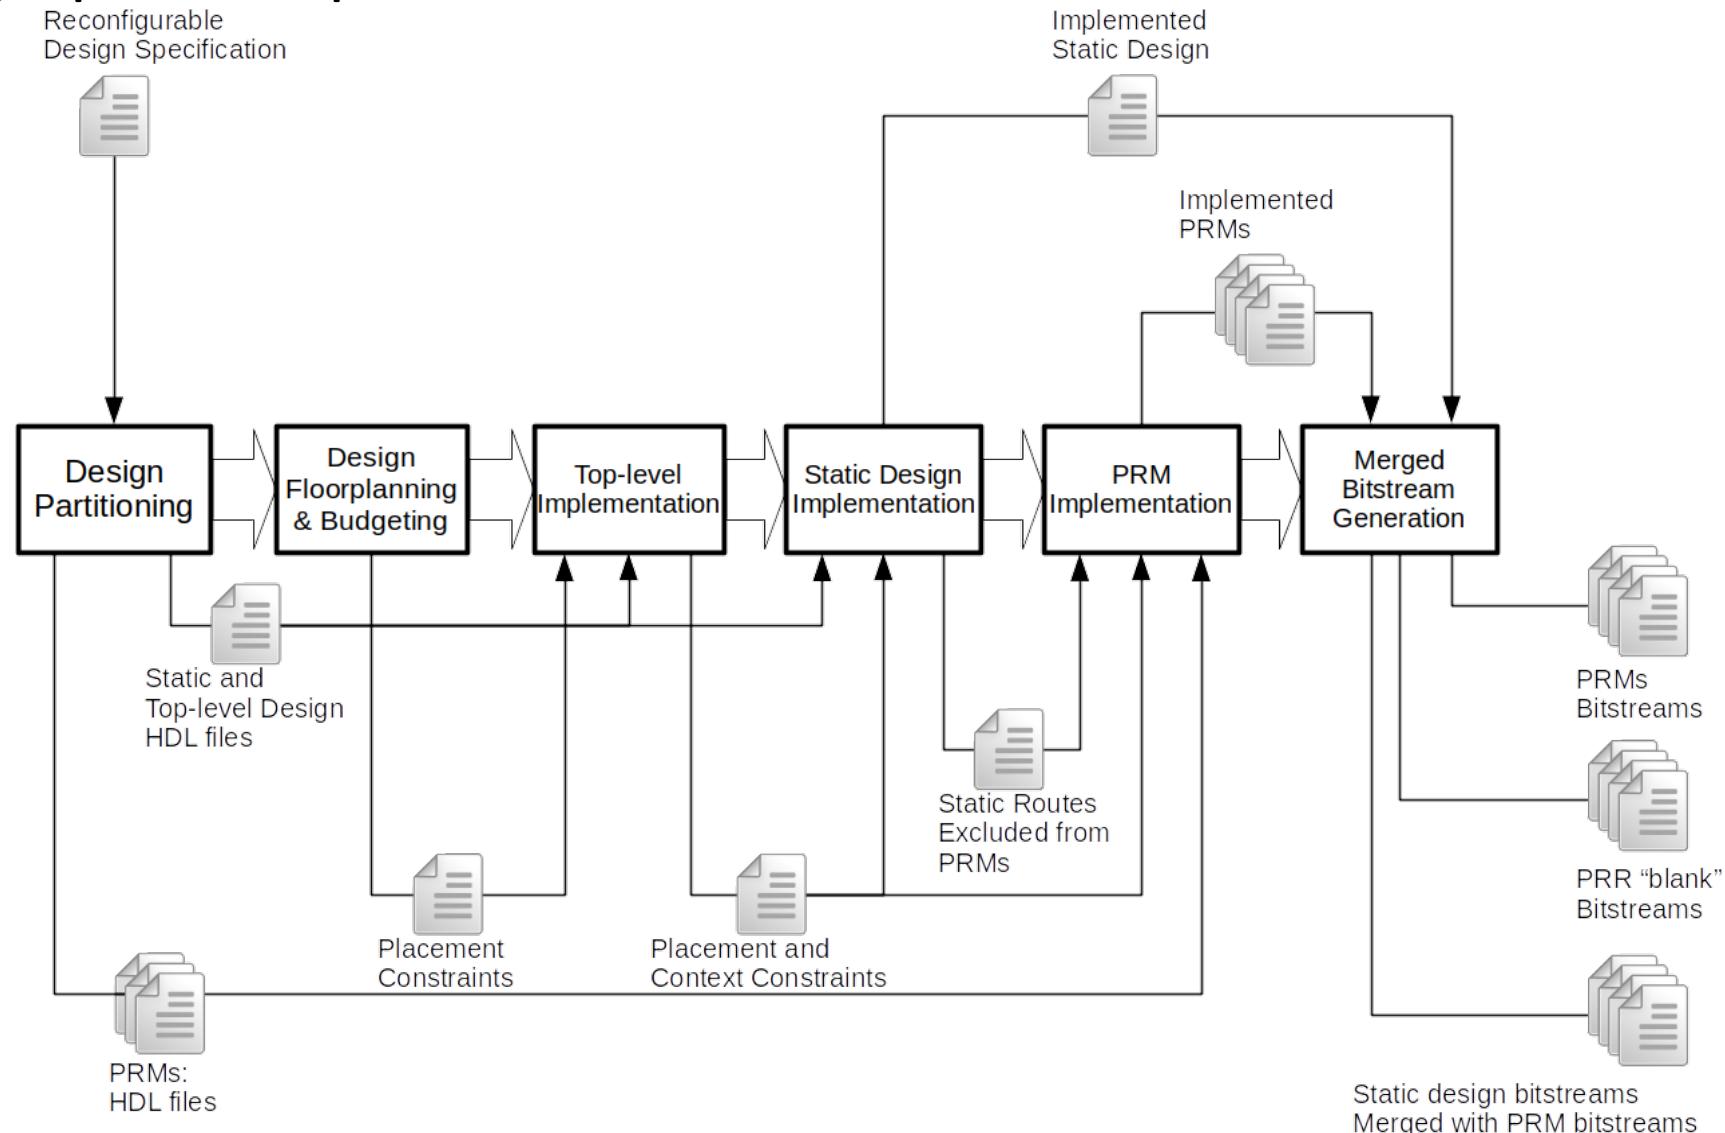
- The size & position of a design cannot be changed
- I/O blocks can only be accessed by contiguous modules
- PRMs/RMs can only communicate with neighbor modules
 - Must be done through bus macros /partition pins
- No global signals allowed
 - E.G., global reset signals are forbidden
 - Exception: clocks which use different bitstreams & routing channels

□ Module-based PR using Vivado

1. Synthesize the static & Reconfigurable Modules separately.
2. Create physical constraints (Pblocks) to define the reconfigurable regions.
3. Set the HD.RECONFIGURABLE property on each Reconfigurable Partition.
4. Implement a complete design (static and one Reconfigurable Module per Reconfigurable Partition) in context.
5. Save a design checkpoint (DCP) for the full routed design.
6. Remove Reconfigurable Modules from this design and save a static-only design checkpoint.
7. Lock the static placement and routing.
8. Add new Reconfigurable Modules to the static design and implement this new configuration, saving a checkpoint for the full routed design.
9. Repeat Step 8 until all Reconfigurable Modules are implemented.
10. Run a verification utility (`pr_verify`) on all configurations.
11. Create bitstreams for each configuration
 - One full bitstream for the static design
 - One partial bitstream for each PRM

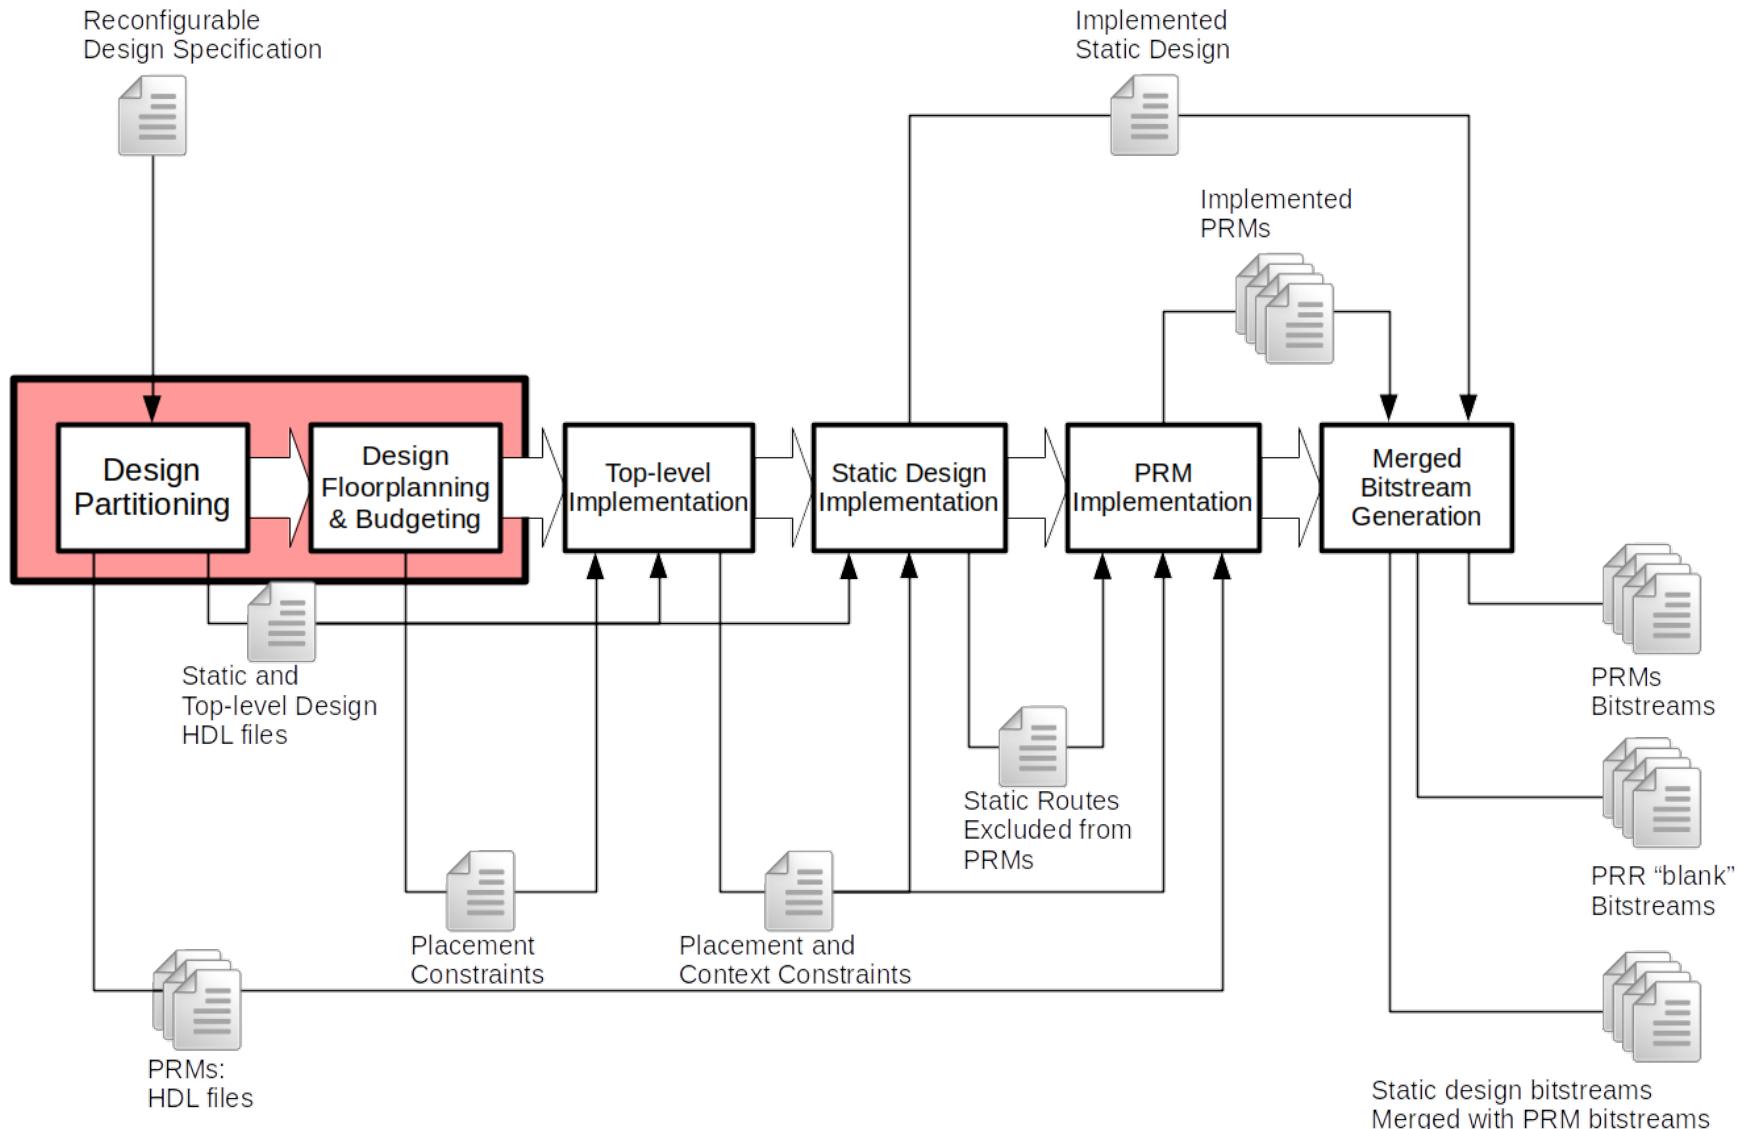
Vivado Software flow

□ A graphical representation of the PR flow



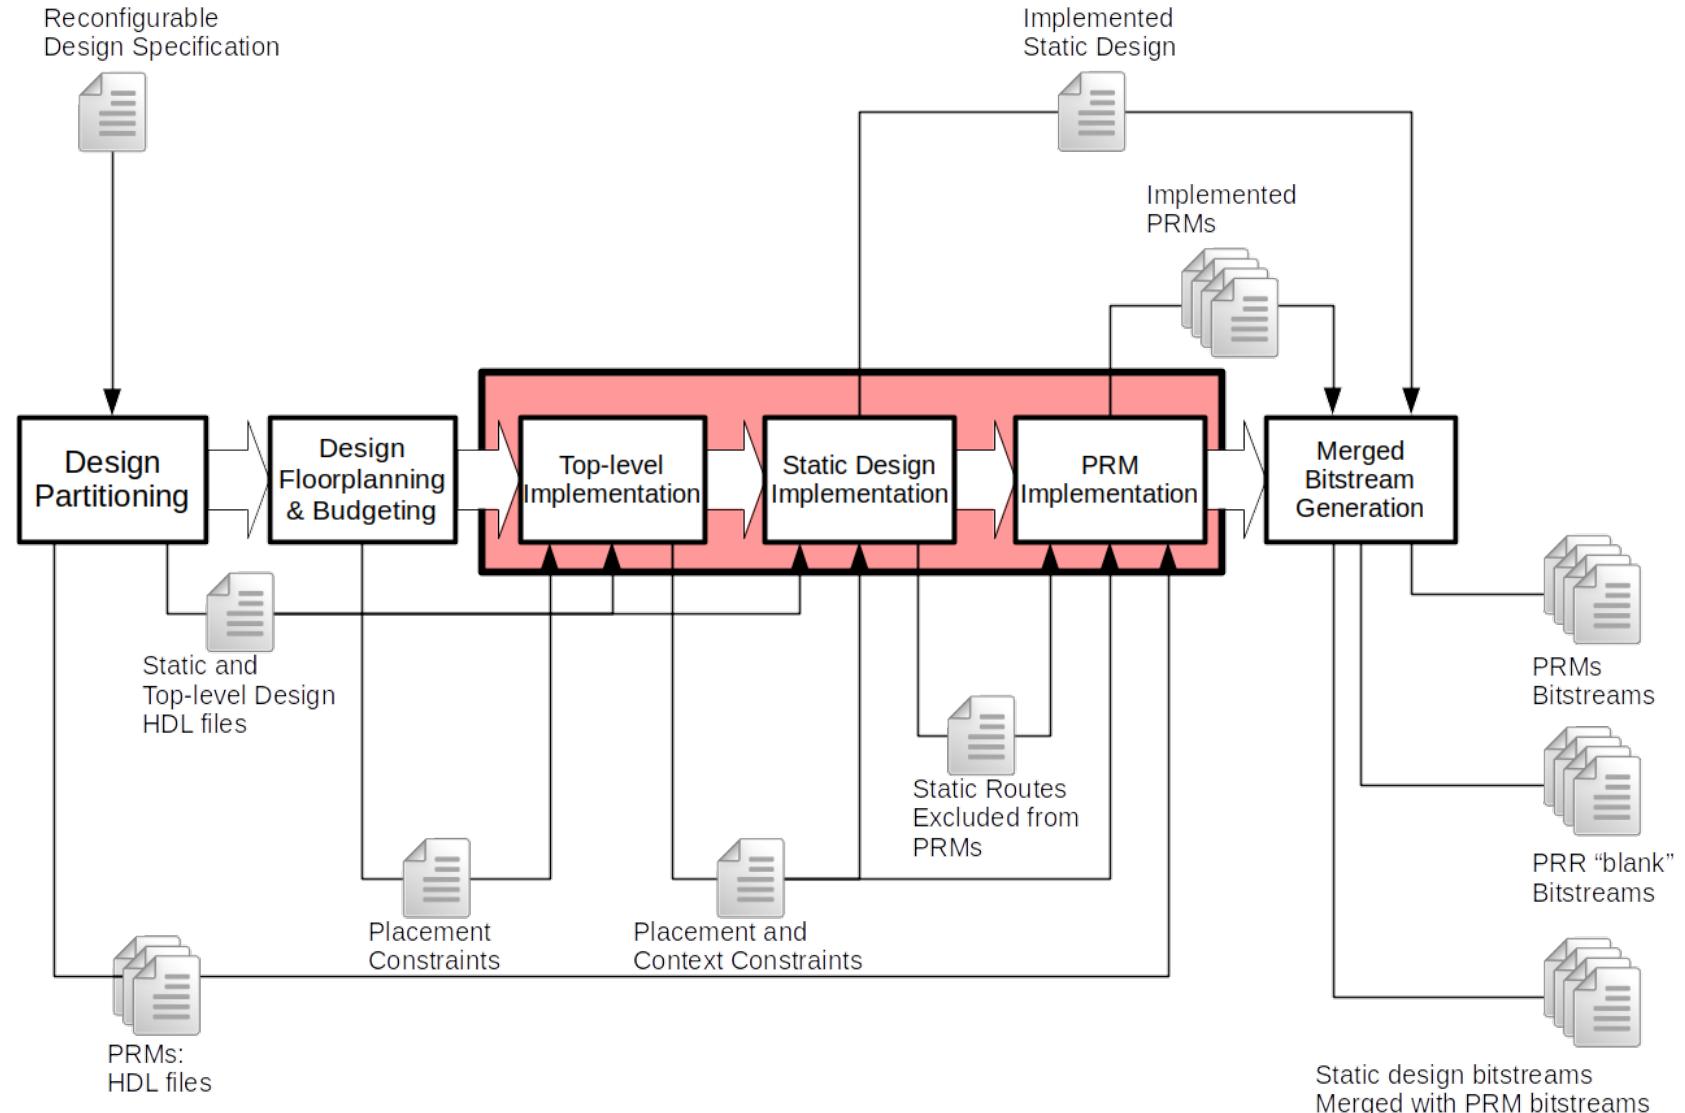
Vivado Software flow

□ A graphical representation of the PR flow



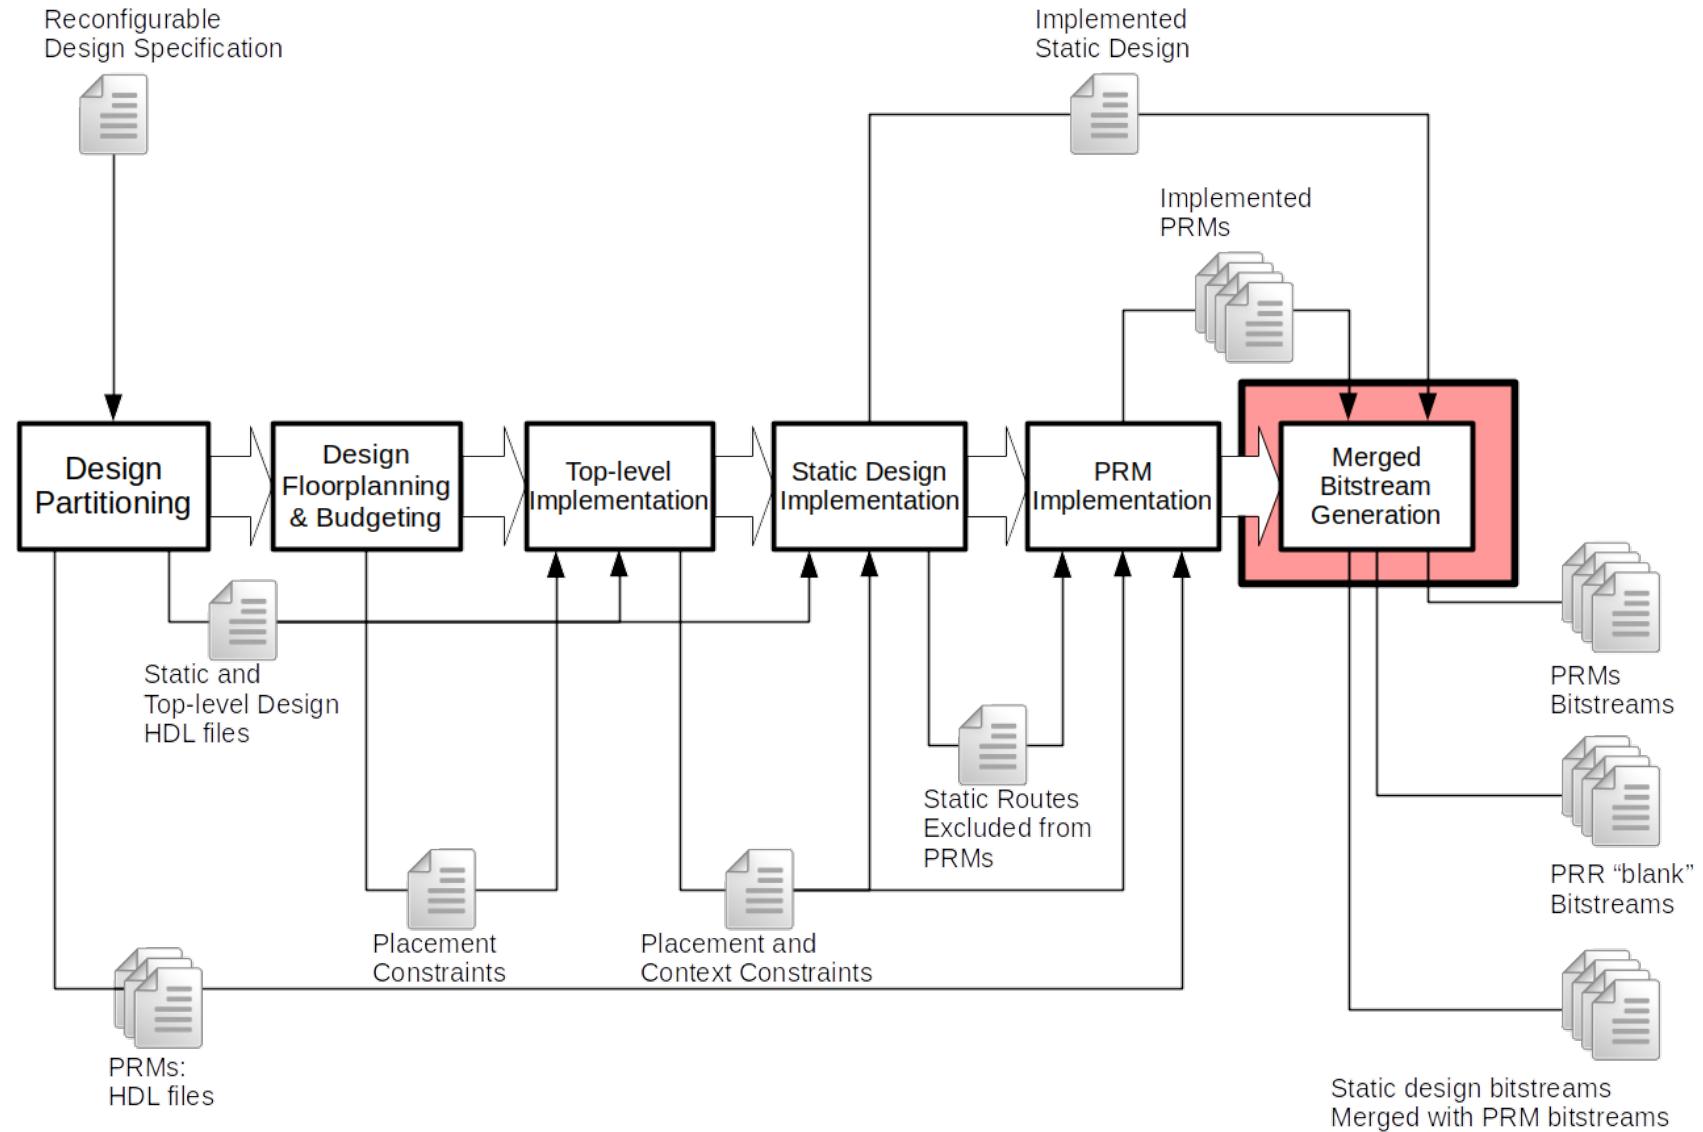
Vivado Software flow

□ A graphical representation of the PR flow



Vivado Software flow

□ A graphical representation of the PR flow



□ Criteria for dynamic PR in FPGAs

- Speeds and methods;
- Design hierarchy limitations w.r.t. PR modules (PRMs)
- Number of allowed PR regions (PRRs)
- Software support to generate PR designs

□ Overview

- Small changes on the FPGA configuration
- Manually done using the FPGA editor
- Relies on full bitstream files
 - The reconfiguration occurs when the new file is selected for reconfiguration
 - A difference is computed between the current and the new file to see what to reconfigure

❑ What can be modified ?

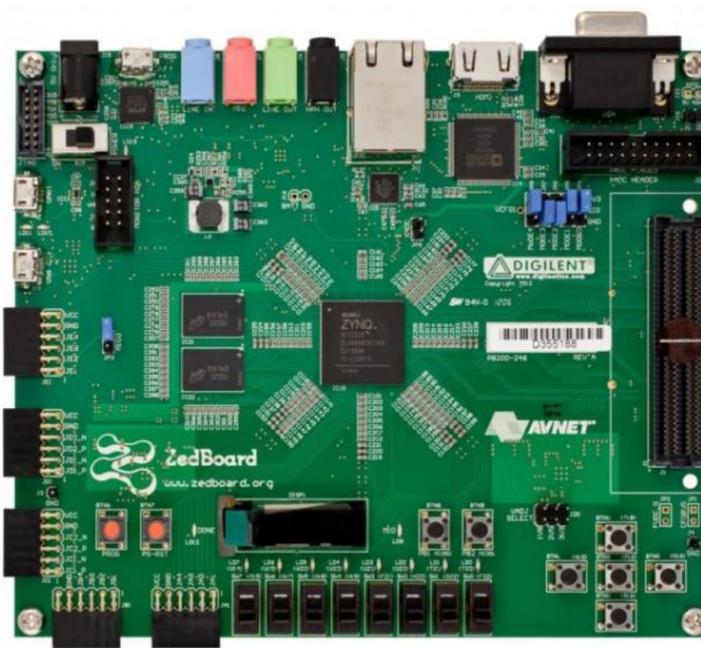
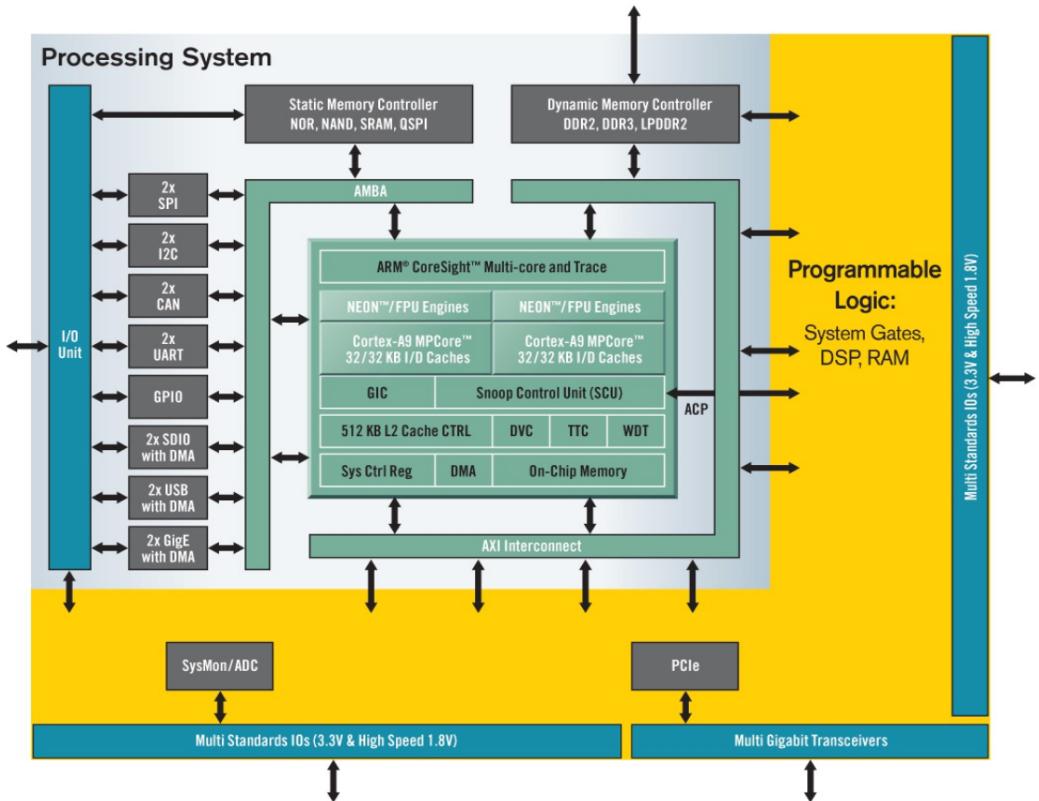
- LUT equations
- BRAM contents and BRAM write modes
- I/O standard pull-ups/pull-downs on external pins
- MUXes to invert polarity
- Flipflop initialization and reset values

❑ What can NOT be modified ?

- Routing—very dangerous: internal contention

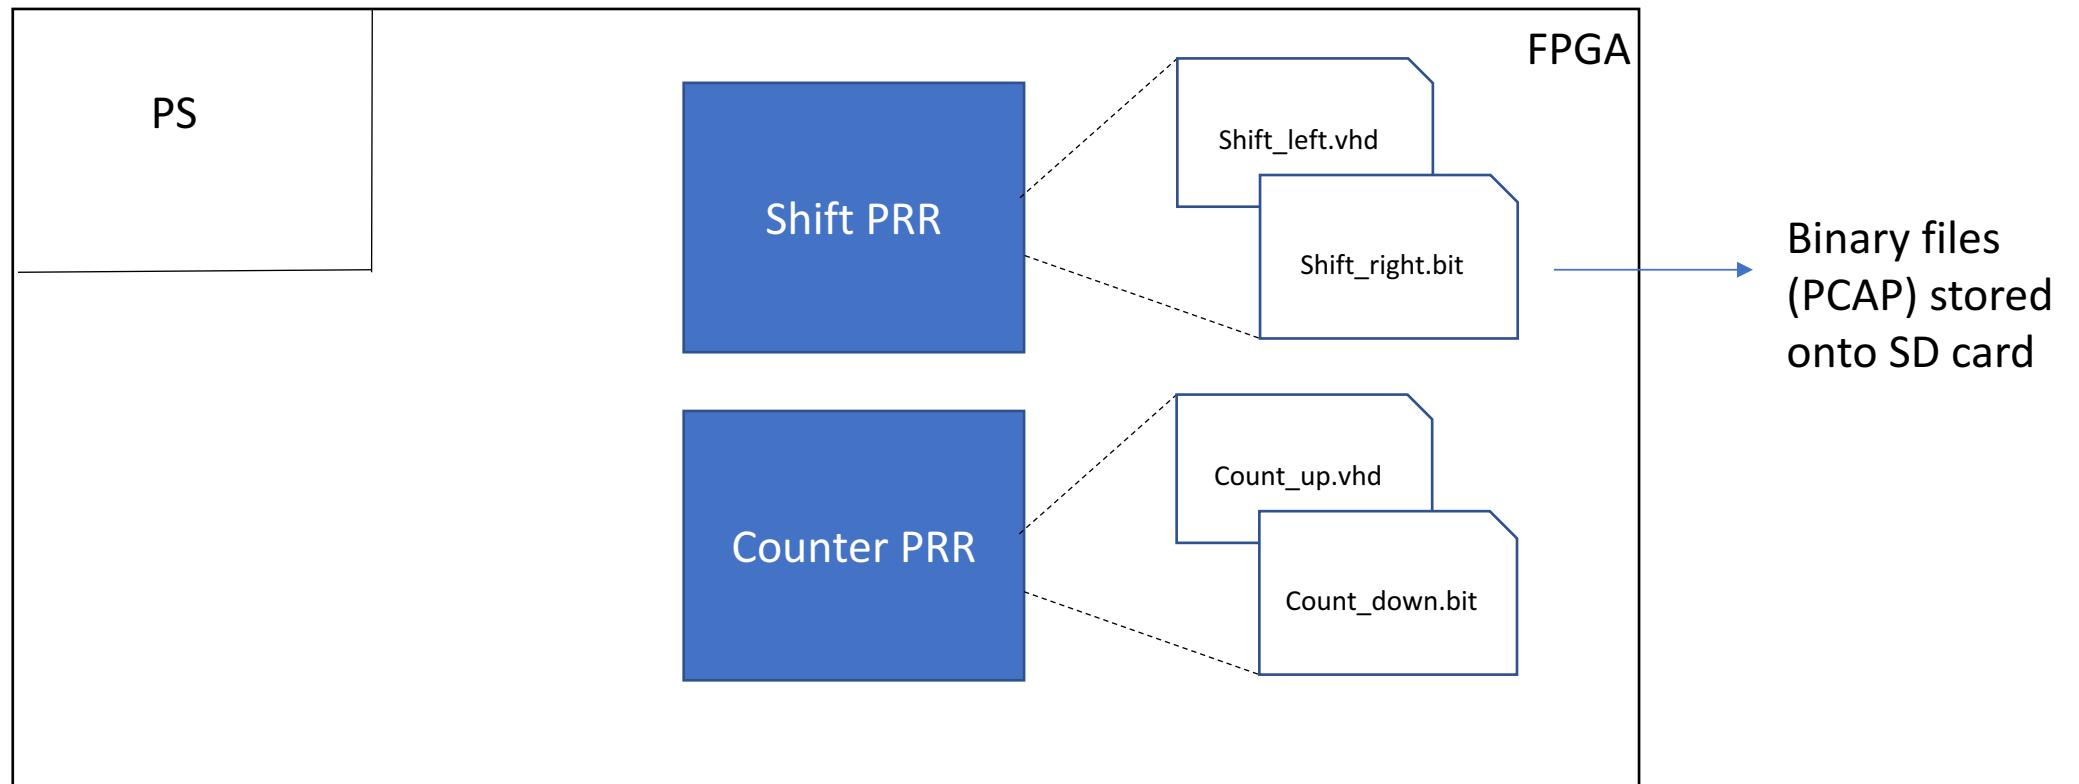
□ During the TP,

- Realization of the PR design flow using Vivado
 - Target platform is zedboard – Zynq-7000
 - Dynamic partial reconfiguration using ARM pro



PRATICAL WORK

- ❑ During the TP, 2 PRRs and 4 bitstreams will be used
- ❑ DPR Management using PCAP and DevCFG
- ❑ Bitstreams (*.bin) onto SD card



Conclusion

- Partial reconfiguration in FPGAs offers a flexible way to adapt a device to specific situations
- PR can be
 - static “stop the world” while the system gets reconfigured, even partially
 - dynamic “only a target portion paused while it is being reconfigured” everything else continues to run;
- Advantages include
 - lower power consumption
 - better runtime adaptation
 - possibility to enhance fault-tolerance
 - etc.
- When considering dynamic PR, the designer must take into account reconfiguration times
 - ⇒ In real-time environments, one must ensure deadlines can be met i.e., reconfiguration time should be included when considering tasks latencies